

rojocOS - An attempt at an efficient OS for cryptography

Romain de Coudenhove, Johan Utterström and Constantin Vaillant-Tenzer

May 15, 2025

Table of Contents

1 Introduction

2 Methods

- Randomness
- Malloc Kernel
- File System

3 Results

4 Conclusion

Making an OS

"No one who isn't already a seasoned developer with years of experience in several languages and environments should even be considering OS Dev yet. A decade of programming, including a few years of low-level coding in assembly language and/or a systems language such as C, is pretty much the minimum necessary to even understand the topic well enough to work in it."
(wiki.osdev.org/Beginner_Mistakes)



Our goal

Modify and improve an WeensyOS system to make it compatible with instant disk encryption (Sabt, Achemlal, and Bouabdallah 2015).

- **Instant Disk Encryption:** Files encrypted with AES using public keys derived from entropy pool.
- **File Deletion:** Quickly and securely remove files
- **Randomness:** Enhanced entropy by a user 16 bits key generation.
- **Secure Memory:** Kernel allocator with zeroing and alignment.

Open source code available on <https://github.com/cvt8/rojocos> (from May 18, 2025).

rojocOS Overview

- **Filesystem:** Tree-based, 1376-byte nodes, AES encryption with public keys, supports file deletion.
- **Randomness:** Entropy from keystrokes and TSC, used for public key generation.
- **Memory Allocator:** Secure kernel heap with zeroing and alignment.
- **User Programs:** Linux-like commands (`ls`, `cat`, `plane`) and shell.

Table of Contents

1 Introduction

2 Methods

- Randomness
- Malloc Kernel
- File System

3 Results

4 Conclusion

Why we upgraded randomness?

- **Old state:** ‘rand()’ was a Linear congruential generator that was always seeded with a fixed constant \Rightarrow identical “random” stream each boot.
- **Our requirements:** unpredictable seeds for safe file storage, potential crypto applications, and still run on QEMU with *no extra hardware*.
- **Solution sketch:**
 - 1 Collect **human entropy** from keyboard timing and rdtsc (CPU cycle count)
 - 2 Keep 16-byte pool in kernel; and blend in with new cycle count randomness.
 - 3 New syscall **sys_getrandom**
 - 4 Backward compatible: kernel and user
- **Key design choice:** simple XOR mixing + feed into LCG algorithm for speed; but note that this is *not* a cryptographically secure random number generator.

Kernel implementation (entropy path)

- `entropy.c/h`

new module

- ▶ `request_user_entropy()`: blocking boot prompt, gathers 16 keystrokes.
- ▶ `refresh_entropy()`: re-prompts after 10 000 outputs (for convenience, should be lower for security).
- ▶ `get_entropy_value()`:
 - Each key byte $b_i = \text{ASCII} \oplus \text{TSC}_{7:0} \oplus \text{TSC}_{15:8} \oplus \text{timing}$
 - i.e. keystroke + low and next-low bits at moment key is read + software counter incremented while polling for key
 - XOR: random if at least one is random
 - Diffusion pass: $b_i \hat{=} b_{(i+7) \bmod 16}$ to avoid local bias
 - On read, select 4 bytes starting at (TSC & 15) and XOR with live TSC.

- `kernel.c`

- ▶ Boot: call `request_user_entropy()` before paging test.

User-space integration

- `lib.c`

- ▶ First call to `rand()` now:

kernel build: `srand(get_entropy_value())` | user build: `srand(sys_getrandom())`

- ▶ LCG step from original OS retained: $x_{n+1} = 1664525 x_n + 1013904223 \pmod{2^{32}}$.

- **Demo program** `p-entropy.c`: prints 32-bit numbers, yields every 50 iterations—verifies stream differs across boots. Asks for new entropy when it is used up (we can set this parameter manually)
- **Build changes**: add `entropy.o` to kernel objects; add `p-entropy` target to process list.

Limitations and future work

Downsides


- Requires human keystrokes \Rightarrow blocks unattended boots / CI.
- 128-bit pool with XOR mixing is not cryptographically secure
- Not safe from attacks. A single `sys getrandom()` call (or any kernel-memory leak) reveals the raw XOR-based pool

Next steps

- Gather extra sources of entropy (disk IRQ, network jitter, hardware randomness).
- Ideally: background entropy daemon; boot proceeds without blocking
Entropy pool continually updates


Overview and Architecture

- **Purpose:** Secure memory allocator for WeensyOS kernel, inspired by iPhone's memory protection.
- **Heap Structure:**
 - ▶ 4KB pages via `page_alloc(PO_KERNEL_HEAP)`.
 - ▶ Blocks with headers (`size`, `next`).
 - ▶ Free list with first-fit strategy.
- **Functions:** `kernel_malloc`, `kernel_free`, `extend_heap`.



Security Features


- **Memory Zeroing:** Allocated and freed memory is zeroed to prevent data leaks.
- **16-Byte Alignment:** Supports encryption (e.g., AES) and SIMD operations.
- **Validation:** Checks sizes and pointers to prevent overflows and double-frees.
- **Encryption Compatibility:** Designed for future cryptographic integration.



Inspired by iPhone's
Secure Enclave

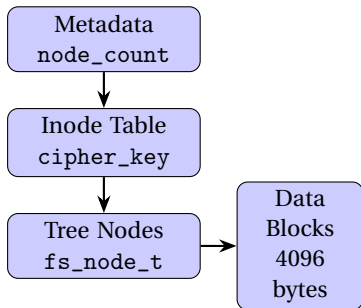
Testing the Implementation

- **Test Function:** testmalloc runs:
 - ▶ Single allocation/free (100 bytes).
 - ▶ Multiple allocations (200, 300 bytes).
 - ▶ Large allocation (2048 bytes).
 - ▶ Stress test (10x 50 bytes).
 - ▶ Zero-size and custom-size tests.
- **Automatic Execution:** Via p-testmalloc on boot.
- **Verification:** Logs in log.txt, memory maps on CGA console.



Key Features of rojocOS Filesystem

- **Tree-Based Organization:** Hierarchical structure rooted at '/'.
- **Fixed-Size Nodes:** `fs_node_t` (1376 bytes), up to 32 children.
- **AES Encryption:** Files encrypted with 256-byte `cipher_key` using public keys from entropy pool.
- **File Deletion:** Securely remove files via `INT_SYS_UNLINK`.
- **Rich System Calls:** `open`, `read`, `write`, `mkdir`, `unlink`, etc.
- **Linux-Like Commands:** `ls`, `touch`, `cat`, `cd`, `mkdir`, `plane`.




Benefits and Design Goals

- **Simplicity:** Minimalist design and small codebase.
- **Efficiency:** Fixed-size nodes (1376 bytes) and 4096-byte blocks optimize disk usage.
- **Security:** AES encryption with public keys (partial implementation via `cipher_key`).
- **Extensibility:** Easy to add new syscalls or user programs.


Filesystem Integration

- **Kernel Initialization:** `fs_init` sets up filesystem with disk callbacks.
- **Process Access:** File descriptors (`proc_fdlist_t`) manage open files.
- **Encryption:** Public keys from entropy pool used for AES key exchange.
- **User Programs:** `cat`, `ls`, etc., use `sys_open()`, `sys_read()`, `sys_write()`.



Filesystem Structure

- **Tree Structure:** Hierarchical, rooted at /.
- **Node (fs_node_t, 1376B):**
 - ▶ **used** (1B): Node status.
 - ▶ **value** (4B): Inode index or 0 (dirs).
 - ▶ **children** (32 entries): 32B name, 4B index each.
- **Storage:**
 - ▶ **tree_usage_offset:** Tracks node availability.
 - ▶ **tree_offset:** Stores nodes.
- **Inodes:** Metadata (start block, block count, 256B AES key).
- **Data Blocks:** 4096B, AES-encrypted.



System Calls

- **INT_SYS_OPEN:** Opens a file, returns file descriptor; uses `fs_getattr`.
- **INT_SYS_READ/WRITE:** Reads/writes encrypted file data; updates descriptor offset.
- **INT_SYS_MKDIR:** Creates directory via `fs_touch (value=0)`.
- **INT_SYS_TOUCH:** Creates file, allocates inode.
- **INT_SYS_UNLINK:** Deletes file, frees inode and blocks.
- **INT_SYS_LISTDIR:** Lists directory contents.
- **INT_SYS_CHDIR/GETCWD:** Manages working directory.

Example: File Operations

```
syscall_open("/file1") → open file.  
syscall_unlink("/file1") → deletes file.
```

User Programs

- **plane**: Text editor; opens files, reads/writes content.
- **shell**: Executes commands via `INT_SYS_EXECV`.
- **mkdir**: Creates directories with `INT_SYS_MKDIR`.
- **ls**: Lists files using `INT_SYS_LISTDIR`.
- **cd**: Changes directory with `INT_SYS_CHDIR`.
- **touch**: Creates files via `INT_SYS_TOUCH`.
- **cat**: Displays file contents using `INT_SYS_READ`.

Shell Example

```
$ ls  
file1 dir1  
$ touch file2  
$ cat file1  
Hello, WeensyOS!
```

Process Management: Fork Syscall

- **INT_SYS_FORK:** Creates a child process:
 - ▶ Allocates new page table for child.
 - ▶ Copies parent's memory pages; shared pages (e.g., kernel) are mapped directly.
 - ▶ Child inherits parent's registers, cwd, but returns 0; parent returns child PID.
- **File Descriptors:** Managed via `proc_fdlst_t`, copied during fork.
- **INT_SYS_EXECV:** Loads new program, preserves arguments.
- **INT_SYS_WAIT/FORGET:** Waits for child exit, frees child resources.




Table of Contents

1 Introduction

2 Methods

- Randomness
- Malloc Kernel
- File System

3 Results

4 Conclusion

Performances tests

Let's make a demo !

```
/$ touch a
/$ plane /a
HELLO
/a opened -> 1
/a writed -> 64
/$ cat /a
read_count : 1024
HELLO/$
/$ mkdir b
mkdir b
/$ ls
a
b
/$ _
```

Table of Contents

1 Introduction

2 Methods

- Randomness
- Malloc Kernel
- File System

3 Results

4 Conclusion

Key Takeaways

- **Secure Filesystem:** Tree-based, AES-encrypted files with public keys, supports deletion.
- **Randomness:** Entropy pool for public key generation, though not cryptographically secure.
- **Memory Allocator:** Secure kernel heap with zeroing and alignment.
- **User Experience:** Linux-like commands and shell for usability.
- **Foundation for Cryptography:** First steps toward a secure OS.

References I



Sabt, Mohamed, Mohammed Achemlal, and Abdelmadjid Bouabdallah (2015). “Trusted execution environment: What it is, and what it is not”. In: *2015 IEEE Trustcom/BigDataSE/Isipa*. Vol. 1. IEEE, pp. 57–64.

Questions



Discussion